

Hadoop Admin Exercise

Killing Hadoop Job

The following command, run in the Hadoop installation directory on a Hadoop cluster, will list all the current jobs:

```
$ bin/hadoop job -list
```

This will produce output that looks something like:

```
1 jobs currently running
JobId   State   StartTime      UserName       Progress
job_200808111901_0001  1       1218506470390  aaron
```

You can use this job id to kill the job; the command is:

```
$ bin/hadoop job -kill jobid
```

Substitute the "job_2008..." from the `-list` command for *jobid*.