

HDFS Administration commands

FS Admin Shell:

HDFS command

Command	Description
bin/hadoop dfs -ls	list files
bin/hadoop dfs -lsr	list in reverse order
bin/hadoop dfs -ls dir1	list files in a directory relative to a user path
bin/hadoop dfs -mkdir somedir	Make a directory
bin/hadoop dfs -put f1.txt somedir	Upload a file
bin/hadoop dfs -put dir1 somedir	Upload a directory
bin/hadoop dfs -cat f1.txt	Display file(s)
bin/hadoop dfs -get f1.txt test.txt	Download a file
bin/hadoop dfs -du dir	Show disk usage
bin/hadoop dfs -dus dir	Show disk usage summary
bin/hadoop dfs -mv f1.txt f2.txt	Move file(s)
bin/hadoop dfs -xp f1.txt f2.txt	Copy file(s)
bin/hadoop dfs -rm f1.txt	Remove file(s)
bin/hadoop dfs -rmr dir1	Remove files recursively
bin/hadoop dfs -moveFromLocal f1.txt f2.txt	Copy and then remove the local copy
bin/hadoop dfs -moveToLocal f1.txt f2.txt	Opposite of moveFromLocal
bin/hadoop dfs -getmerge src local	Download the src in HDFS and merged all files into one
bin/hadoop dfs -setrep 3 path	Change the replication factor
bin/hadoop dfs -touch path	Touch file(s)
bin/hadoop dfs -test path	Check if exists
bin/hadoop dfs -stat path	Print path info
bin/hadoop dfs -tail -f file	tail a file
bin/hadoop dfs -chmod -R mode path	Change a file access right
bin/hadoop dfs -chgrp -R group path	Change a group

DFS Admin Shell

'bin/hadoop dfsadmin' command supports a few HDFS administration related operations. bin/hadoop dfsadmin -help lists all the commands currently supported

bin/hadoop dfsadmin -report	Overall health of HDFS
bin/hadoop dfsadmin -metasave file1	Save the meta data of NameNode in file1

HDFS Admin command

Command	Description
-report	Reports HDFS statistics
-safemode	Enter or leave Safemode
-finalizeUpgrade	Removes previous backup of the cluster
-refreshNodes	Updates the set of hosts allowed to connect to namenode
-printTopology	Print the topology of the cluster

Misc commands:

Fsck

HDFS supports `fsck` command to check for various inconsistencies. It is designed for reporting problems with various files, for e.g. missing blocks for a file or under replicated blocks. Unlike a traditional fsck utility for native filesystems, this command does not correct the errors it detects. Normally Namenode automatically corrects most of the recoverable failures. HDFS' fsck is not a Hadoop shell command

```
'bin/hadoop fsck'
```

Upgrade and Rollback

Before upgrading Hadoop software, *finalize* if there an existing backup.

```
'bin/hadoop dfsadmin -upgradeProgress status'
```

More details : http://wiki.apache.org/hadoop/Hadoop_Upgrade